

Chapter 1

Regression Diagnostics

1.1 Coefficient of determination

1.1.1 R^2 statistic

$$FVU(f) = \frac{MSE(f)}{Var(y)} \leftarrow \text{fraction of variance unexplained}$$

$$R^2 = 1 - FVU(f) = 1 - \frac{MSE(f)}{Var(y)}$$

$$R^2 = 0 \rightarrow \text{Trivial predictor}$$

$$R^2 = 1 \rightarrow \text{Perfect predictor}$$

1.1.2 Overfitting

When a model performs well on training data but doesn't generalize, we are said to be overfitting.

$$X\theta = y, \quad \theta \leftarrow \text{hypothesis}$$

A "simple" model is one where θ

- has few non-zero parameters
- is almost uniform

1.1.3 Regularization

Regularization is the process of penalizing model complexity during training

$$\arg \min \theta = \frac{1}{N} \|y - X\theta\|_2^2 + \lambda \|\theta\|_2^2$$

Optimizing the (regularized) model

$$\frac{\partial f}{\partial \theta_k} = \frac{1}{N} \sum_i -2x_{ik}(y_i - x_i\theta) + 2\lambda\theta$$

1.1.4 Model Selection

How to select which model is best? We need a **third** dataset.

A **validation set** is constructed to "tune" the model's parameters that are not directly optimized. Some

- The training error increases as λ increases
- The validation and test error are at least as large as the training error
- The validation/test error will usually have a "sweet spot" between under- and over-fitting

1.2 Summary

1. Regression can be cast in terms of maximizing a likelihood
2. Gradient descent for model optimization
 - Initialize θ at random
 - While (not converge) do $\theta := \theta - \alpha f'(x)$
3. Regularization is the process of penalizing model complexity during training
4. Regularization Pipeline
 - Training
 - Test
 - Validation

Chapter 2

Supervised learning - Classification

2.1 Naive Bayes

Naive Bayes assumes that features are conditionally independent given the label

$$P(\text{label}|\text{features}) = \frac{P(\text{label}) \prod_i P(\text{features}_i|\text{label})}{P(\text{features})}$$

The denominator doesn't matter, because we really just care about

$$P(\text{label}|\text{features}) \quad \text{V.S.} \quad P(\text{not label}|\text{features})$$

2.2 Logistic Regression

sigmoid function: $\sigma(t) = \frac{1}{1+e^{-t}}$. **Training:**

$$\begin{aligned} L_\theta(y|X) &= \prod_{y_i=1} P_\theta(y_i|X_i) \prod_{y_i=0} (1 - P_\theta(y_i|X_i)) \\ \log L_\theta(y|X) &= \sum_i -\log(1 + e^{-X_i \cdot \theta}) + \sum_{y_i=0} -X_i \cdot \theta - \lambda \|\theta\|_2^2 \\ \frac{\partial \log L_\theta(y|X)}{\partial \theta_k} &= \sum_i \frac{x_{ik} e^{-X_i \cdot \theta}}{1 + e^{-X_i \cdot \theta}} + \sum_{y_i=0} -x_{ik} - 2\lambda \theta_k \\ &= \sum_i x_{ik} [1 - \sigma(X_i \cdot \theta)] + \sum_{y_i=0} -x_{ik} - 2\lambda \theta_k \end{aligned}$$

2.2.1 Multiclass classification

The most common way to generalize binary classification (output in $\{0,1\}$) to multiclass classification (output in $\{1 \dots N\}$) is simply to a binary predictor for each class.

2.3 Support Vector Machines

Distance from a point to a line

$$\begin{aligned}
 ax + by + c &= 0 \\
 d(\text{line}, p_0) &= \frac{|ax_0 + by_0 + c|}{\sqrt{a^2 + b^2}} \\
 \theta x - \alpha &= 0 \Rightarrow \frac{|\theta x - \alpha|}{\|\theta\|_2} \\
 \frac{|\theta x - \alpha|}{\|\theta\|_2} y &\geq 1
 \end{aligned}$$

Want the margin $\frac{1}{\|\theta\|}$ to be as wide as possible, while penalizing points on the wrong side. Soft-margin formulation:

$$\arg \min_{\theta, \alpha, \epsilon_i > 0} \frac{1}{2} \|\theta\|_2^2 + \epsilon$$

Try to optimize the **misclassification error** rather than maximize a probability.

2.4 Evaluate Classifiers

classification accuracy = correct prediction / #predictions = (TP+TN) / (TP+TN+FP+FN)

Error rate = incorrect prediction / #predictions = (FP+FN) / (TP+TN+FP+FN)

True positive rate (TPR) = TP / (TP + FN)

True negative rate (TNR) = TN / (TN + FP)

Balanced Error Rate (BER) = 1/2 (FPR + FNR) = 1 - 0.5*(TPR + TNR)

= 1/2 FOR A RANDOM/NAIVE classifier, 0 for a perfect classifier.

How to optimize a balanced error measure:

$$L_\theta(y|X) = \prod_{y_i=1} p_\theta(y_i|X_i) \prod_{y_i=0} (1 - p_\theta(y_i|X_i))$$

Then,

$$\begin{aligned}
 l_\theta(y|X) &= \sum_{y_i=1} \log \sigma(X_i \cdot \theta) + \sum_{y_i=0} \log(1 - \sigma(X_i \cdot \theta)) \\
 &= \frac{1}{|y_i = 1|} \sum_{y_i=1} \log \sigma(X_i \cdot \theta) + \frac{1}{|y_i = 0|} \sum_{y_i=0} \log(1 - \sigma(X_i \cdot \theta))
 \end{aligned}$$

2.4.1 ranking

The classifiers we've seen can associate scores with each prediction

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

(harmonic mean of precision and recall)

$$F_\beta = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{\beta^2 \cdot \text{precision} + \text{recall}}$$

(weighted in case precision is more important (low beta), or recall is more important (high beta))

2.5 Case study

2.5.1 Using Regression to Predict Content Popularity of Reddit

Chapter 3

Dimensionality Reduction

A1: A (sparse) vector including all movies Raw data issues:

1. Incredibly high-dimensional
2. Missing values
3. Not stable/robust to changing data
4. Redundant dimensions

A2: Describe my preferences using a **low-dimensional** vector. Represent each node/user in terms of the communities they belong to

Goal: take **high-dimensional** data, and describe it compactly using a small number of dimensions.

Assumption: Data lies (approximately) on some **low-dimensional** space.

Why dimensionality reduction? Unsupervised learning

To understand the important features of a dataset and the process which generated the data itself. The models we learn will prove useful when it comes to solving predictive tasks later.

3.1 Principal Component Analysis

- To select a few important features
- To compress the data by ignoring components which aren't meaningful.

For a single data point: $y = \underbrace{\varphi x}_{\text{compress}}, x = \underbrace{\varphi^{-1}y = \varphi^T y}_{\text{decompress}}$

$$x = \varphi_1 y_1 + \varphi_2 y_2 + \dots + \varphi_M y_M = \underbrace{\sum_{j=1}^M \varphi_j y_j}_{\text{"complete" reconstruction}}$$
$$x_i \simeq \underbrace{\sum_{j=1}^k \varphi_j y_j + \sum_{j=k+1}^M \varphi_j b_j}_{\text{approximate reconstruction}}$$

It should minimize the **MSE**:

$$\begin{aligned} \min_{\varphi, b} \frac{1}{N} & \left\| \sum_{j=1}^k \varphi_j y_j + \sum_{j=k+1}^M \varphi_j b_j - \varphi^T y \right\|_2^2 \\ &= \frac{1}{N} \sum_y \left\| \sum_{j=1}^M \varphi_j (y_j - b_j) \right\|_2^2 \\ &= \sum_y \sum_{j=k+1}^M (y_j - b_j)^2 \end{aligned}$$

where $b_j = \bar{y}_j$

MSE is equal to the variance in the discarded dimensions.

Expand in terms of X

$$\begin{aligned} \sum_{j=k+1}^M \sum_i [\varphi_j (X - \bar{X})]^2 &= \frac{1}{N} \sum_{j=k+1}^M \varphi_j (X - \bar{X})^T (X - \bar{X}) \varphi_j^T \\ &= \frac{1}{N} \sum_{j=k+1}^M \varphi_j \text{cov}(X) \varphi_j^T \end{aligned}$$

Solve:

$$\frac{\partial}{\partial \varphi_j} \frac{1}{N} \sum_{j=k+1}^M \varphi_j \text{cov}(X) \varphi_j^T - \lambda_j (\varphi_j \varphi_j^T - 1) = 0$$

This expression can only be satisfied if φ_j and λ_j are an eigenvectors/eigenvalues of the covariance matrix.

3.2 K-means Clustering

1. Input is still a matrix of features
2. Output is a list of cluster "centroids"
3. From this we can describe each point in X by its cluster membership

Given feature(X) our goal is to choose k centroids (C) and cluster assignments (Y) so that the reconstruction error is minimized.

$$\text{reconstruction error} = \sum_i \|X_i - C_{y_i}\|_2^2$$

This is an NP-Hard optimization problem. Use Greedy algorithm:

Initializing C (e.g. at random)

Do

Assign each X_i to its nearest centroid

Update each centroid to be the **mean** of points assigned to it

While (assignments change between iterations)